


PiCAS: New Design of Priority-Driven Chain-Aware Scheduling for ROS2

Hyunjong Choi, Yecheng Xiang, Hyoseung Kim


Robot Operating System (ROS)

- ROS (since 2007)
 - Popular open-source middleware in academia and industry
 - Provides software tools, robot systems, and best-practices

➔ Over the decades, it has revealed shortcomings in real-time support for timing- and safety-critical applications



ROS 1.0
<http://ros.org/>

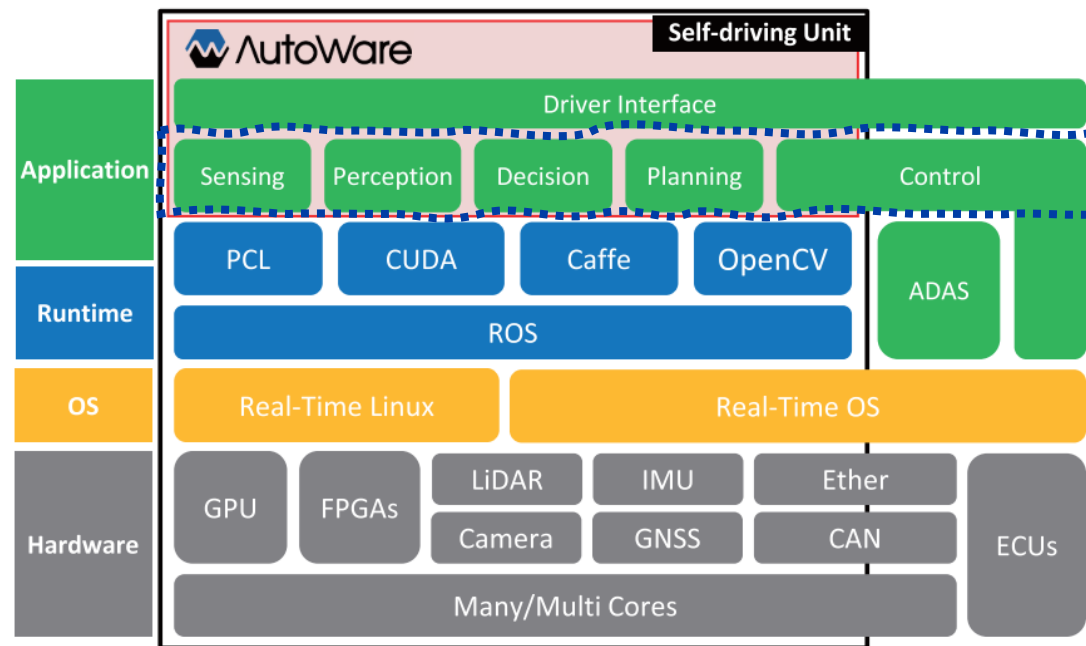


Willow Garage PR2
(original ROS robot)

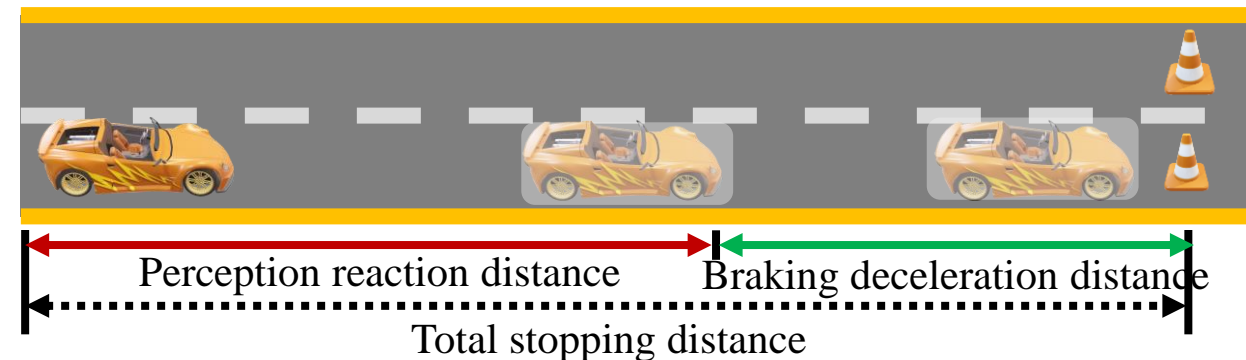
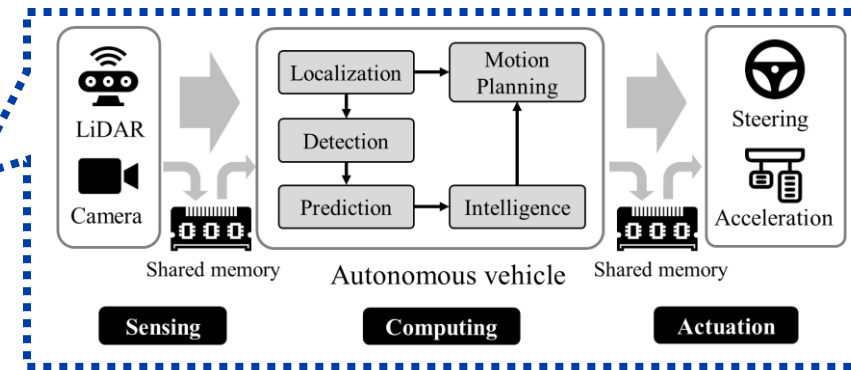
<http://willowgarage.com>

Why real-time in ROS ?

- To develop safety-critical application with ROS
 - Autonomous driving software (e.g., autoware.ai)



< Autoware.ai > [†]



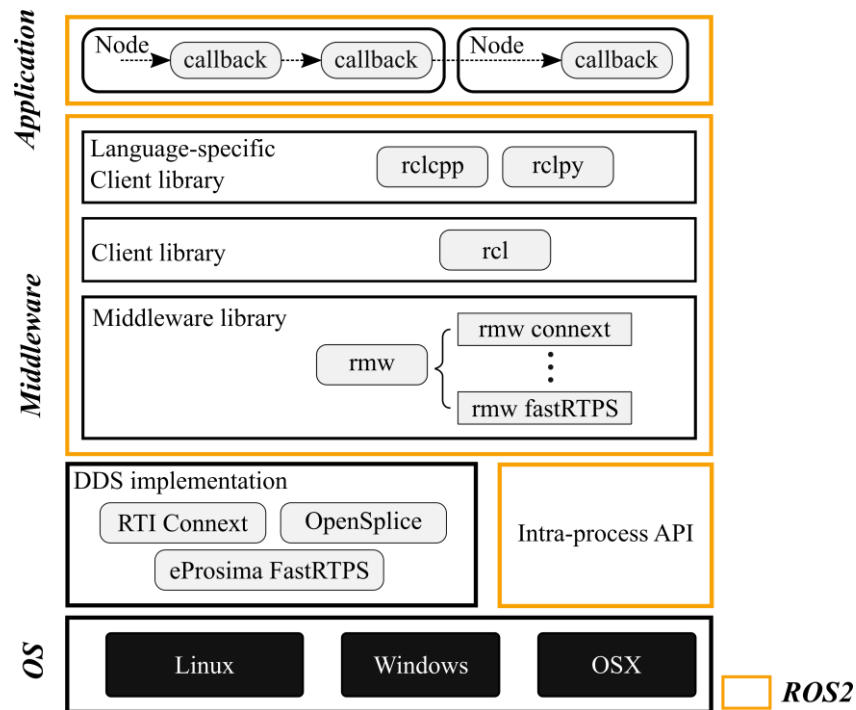
< Motivational example: chain in self-driving application >

➡ Timing constraint violations (e.g., end-to-end latency) can cause catastrophic accidents

[†]S. Kato et al. "Autoware on Board: Enabling Autonomous Vehicles with Embedded Systems", ICCPS, 2018

ROS 2 (since 2017)

- Most concepts are inherited from the original ROS design (e.g., pub-sub)
- Aims to improve **real-time capability, QoS, and security**
- Supports Data Distribution Service (DDS)



< ROS2 architecture >



Ardent Apalone,
released Dec 2017



Eloquent Elusor,
released Nov 2019



Galactic Geochelone,
released May 2021

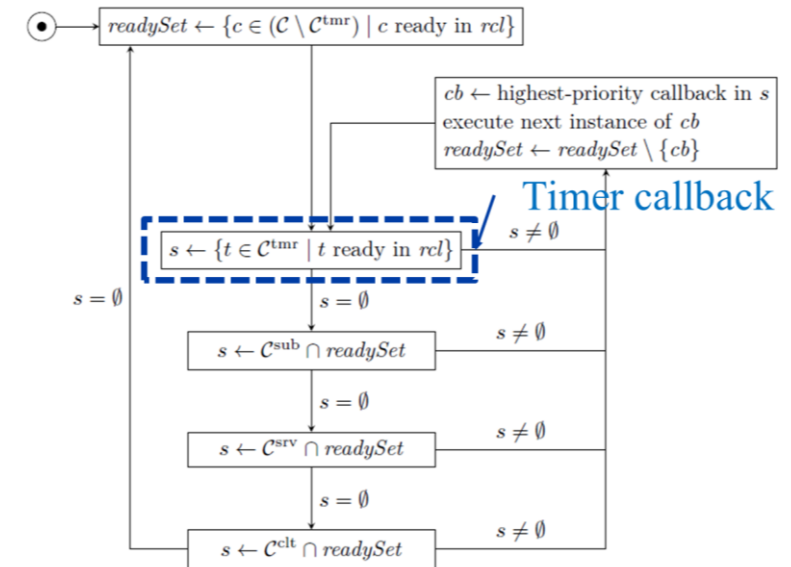
Limitation of current ROS2

- Priority-unaware complex layers of abstractions
 - Unique callback scheduling behavior
 - Ignores criticality or urgency of task chains*

➡ Suffers from priority inversion

- Lack of systematic support for resource allocation and analysis
 - Poor resource utilization
 - Nondeterministic end-to-end behavior

➡ Needs a new RT scheduling framework for ROS2 !



< Callback scheduling in executor[†] >

[†] D. Casini et al. "Response-time analysis of ROS 2 processing chains under reservation-based scheduling", ECRTS, 2019

Scheduling-related abstractions in ROS2


▪ Callbacks, nodes, and executors

Callback model

- ✓ Timer / regular callbacks
- ✓ Non-preemptive
- ✓ $\tau_i := (C_i, D_i, T_i, \pi_i)$

Chain model

- ✓ $\Gamma^c := [\tau_s, \tau_{m1}, \dots, \tau_e]$
- ✓ τ_s : the start callback of Γ^c
- ✓ τ_{m*} : the intermediate callback of Γ^c
- ✓ τ_e : the end callback of Γ^c



Node model

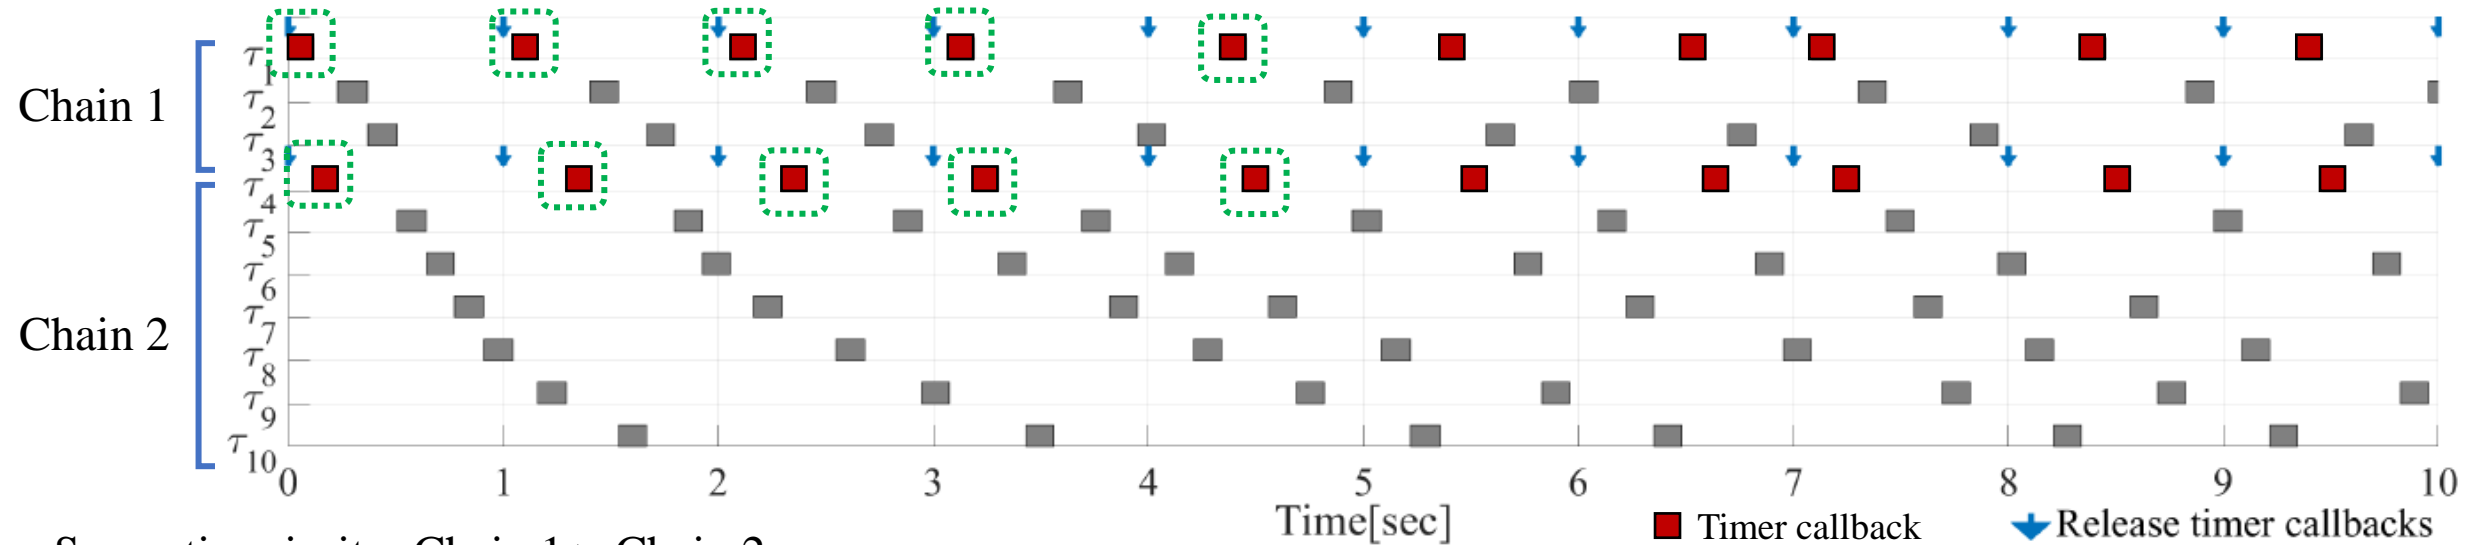
- ✓ $\mathcal{N} =: \{n_1, \dots, n_j, \dots, n_N\}$
- ✓ Not schedulable entities

Executor model

- ✓ $\mathcal{E}_i =: \{e_1, \dots, e_j, \dots, e_E\}$
- ✓ Preemptive
- ✓ Schedule with SCHED_FIFO

Challenges (1/2)

- Challenge I: Fairness-based scheduling *within executors*



Semantic priority: Chain 1 > Chain 2

Single executor	Mean	Max	Min	STD
Chain 1	36.865	72.752	0.505	21.223
Chain 2	36.730	73.149	0.773	21.154

< End-to-end latency results [sec] >

- O1. Prioritizes timer callbacks regardless of chain priority[†]
- O2. Does not distinguish callbacks by their origin chains

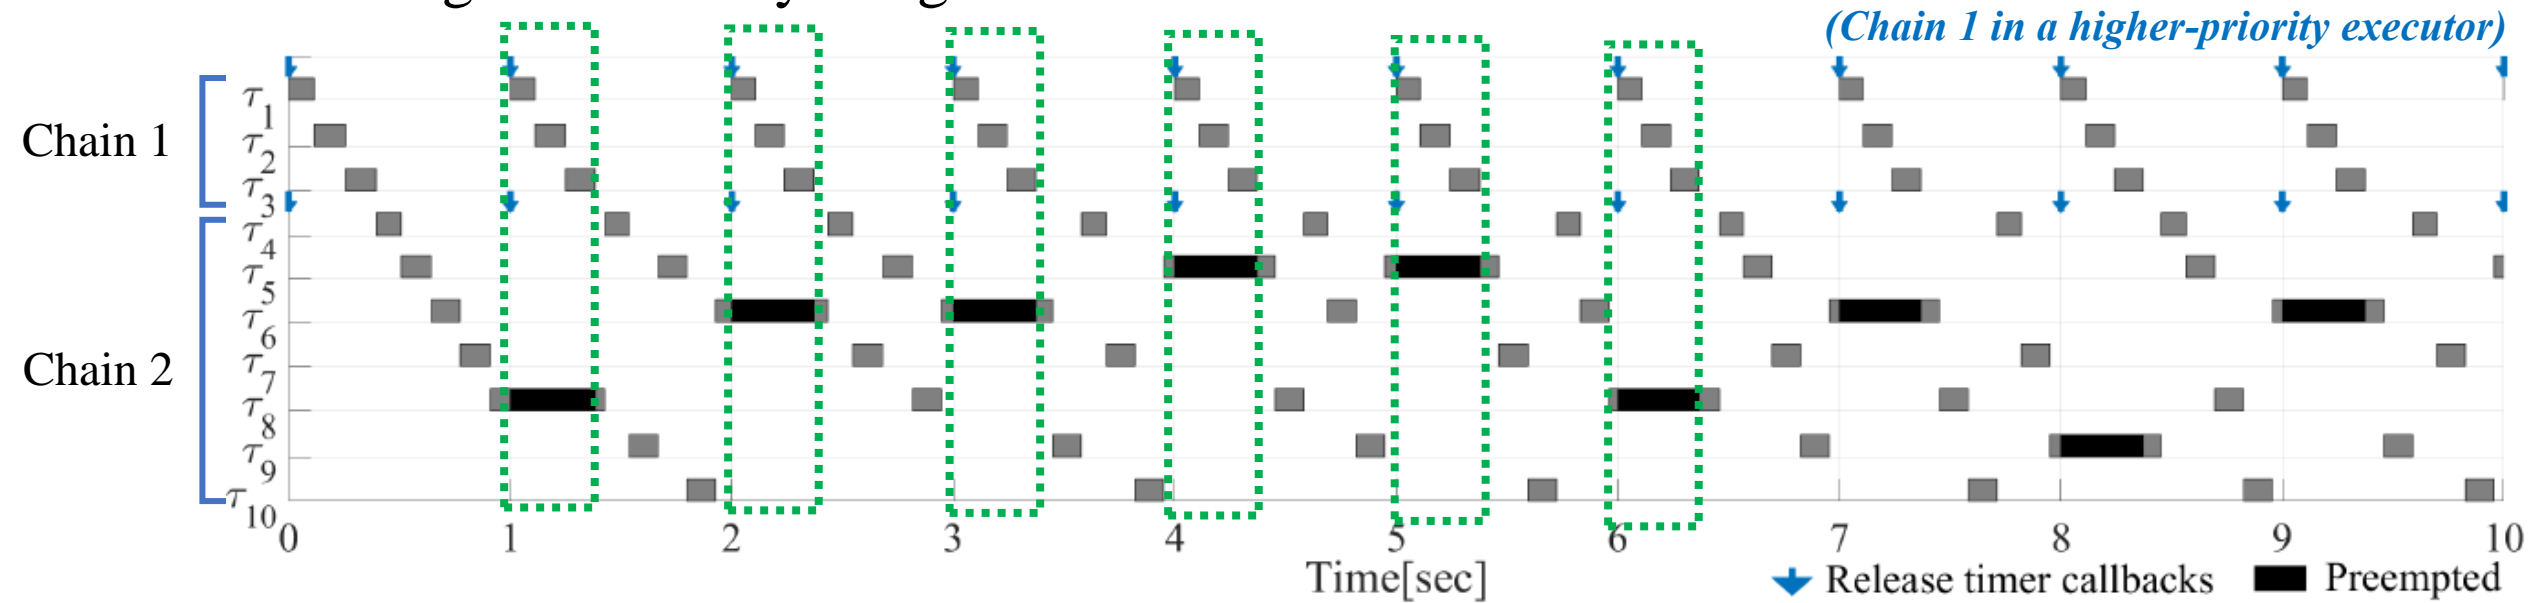


Jeopardizes the timeliness of safety-critical chains

[†] D. Casini et al. "Response-time analysis of ROS 2 processing chains under reservation-based scheduling", ECRTS, 2019

Challenges (2/2)

- Challenge II : Priority assignment for executors



Semantic priority: Chain 1 > Chain 2

Single executor	Mean	Max	Min	STD
Chain 1	0.370	0.392	0.366	0.004
Chain 2	48.795	97.783	0.772	28.304

< End-to-end latency results [sec] >

O3. High penalty due to self-interference

O4. No guidelines on executor priority assignment




Default ROS2 causes unacceptably high latency for chain 2

PiCAS: new real-time scheduling and analysis framework for ROS2

- Key idea: enables *prioritization of critical computation chains* across complex abstraction layers of ROS2
 - To minimize end-to-end latency
 - To ensure predictability even when the system is overloaded
- Re-design ROS2 scheduling architecture to achieve
 - (1) *Higher-semantic priority chain* executes *first*.
 - (2) *For each chain*, its instances on the same CPU *execute in arrival order to prevent self-interference*.

Scheduling strategies

- Strategies for chains running within an executor



	Regular callbacks only	Timer and regular callbacks
Single chain	Strategy I. (To satisfy ① of Lemma 1) 	Strategy II. (To satisfy ① of Lemma 1)
Multiple chains	Strategy III. 	Strategy IV.

- Strategies for chains running across executors

Single chain on one CPU	Multiple chains on one CPU
Strategy V. (To satisfy ② of Lemma 1) 	Strategy VI.

Priority assignment

- *Realization of scheduling strategies* in two aspects
 - Callback priority assignment
 - Chain-aware node allocation algorithm

Algorithm 1 Callback priority assignment

Input: Γ : chains

```

1:  $\Gamma \leftarrow$  sort in ascending order of semantic priority  $\pi_\Gamma$ 
2:  $p \leftarrow 1$  ▷ Initialize current priority
3: for all  $\Gamma^c \in \Gamma$  do
4:   for all  $\tau_i \in \Gamma^c$  do
5:      $\tau_i \leftarrow p$ 
6:      $p \leftarrow p + 1$ 
7:   end for
8: end for
  
```

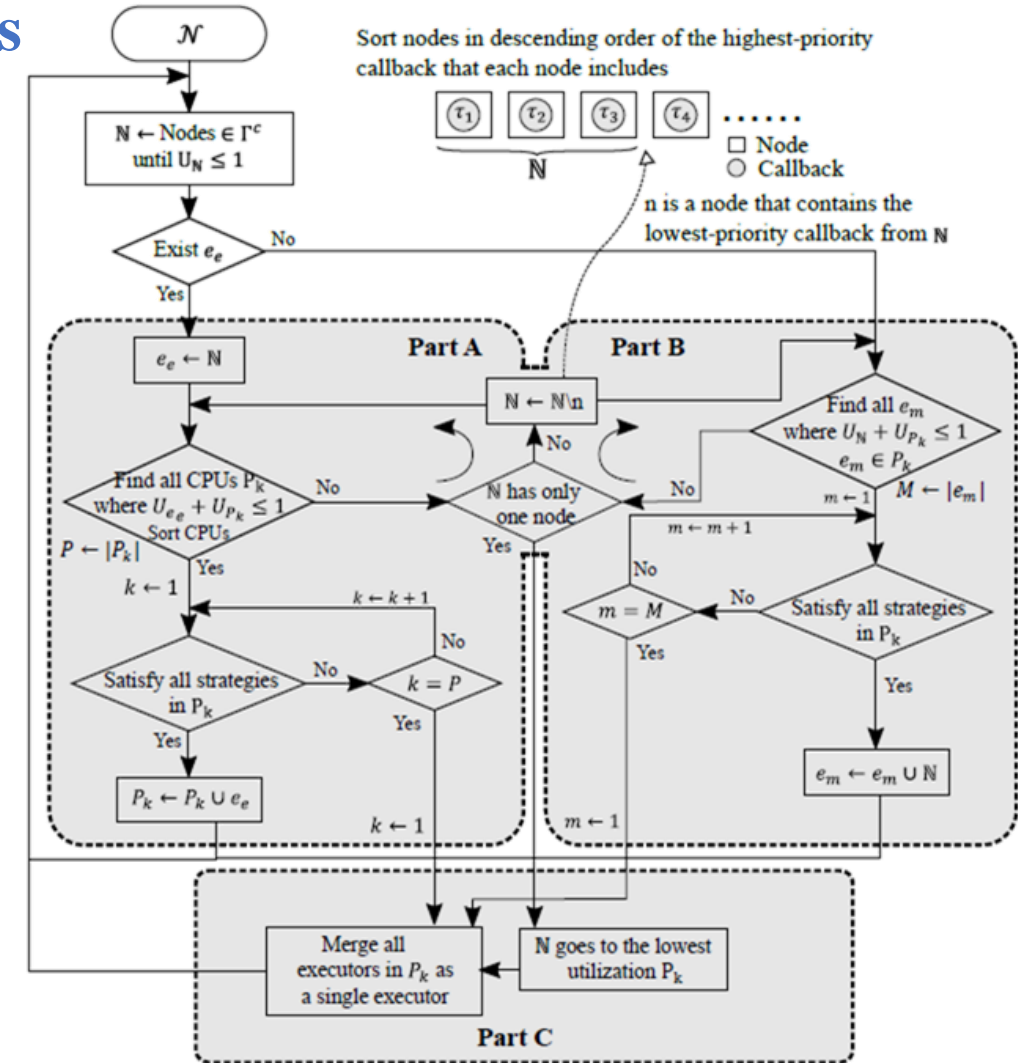
Chain-aware node allocation

- Purpose: **minimize interference between chains**
 - (1) allocate given *nodes to executors*, and then
 - (2) maps *executors to available CPU cores*

* **Allocate all nodes associated to one chain to the same CPU core whenever it is possible**


➔ Reduce end-to-end latency of critical chains!

	Parameters		
\mathcal{N}	Nodes	e_m	Non-empty executors
\mathbb{N}	A node set consists of callbacks of a chain Γ^c ($U_{\mathbb{N}} \leq 1$)	M	The number of e_m
e_e	Empty executor	P	The number of P_k
U_{P_k}	Utilization of CPU core P_k		
n	A node that has the lowest priority callback of Γ^c in \mathbb{N}		




Examples of chain-aware scheduling

- With the same workload at challenges slide.



< All callbacks in a single executor >



< One executor per chain >

Single executor	Mean	Max	Min	STD
Chain 1	0.436	0.506	0.368	0.038
Chain 2	1.196	1.738	0.741	0.348

Executor per chain	Mean	Max	Min	STD
Chain 1	0.369	0.394	0.366	0.004
Chain 2	1.255	1.731	0.737	0.352



Significantly improved end-to-end latency under PiCAS

Analysis of end-to-end latency

- Latency analysis in a multi-core system
 - Segment Φ_i : a subset of a chain on one CPU core
 - Multiple segments if a chain executes over multiple CPU cores

Step 1: Computing the WCRT of each segment of a chain

WCRT of a segment $\Phi_i, R_{c,i}^n$



Step 2: Adding the WCRT of all segments of the chain

End-to-end latency of a chain, L_{Γ^c}

< Latency analysis of a chain in a multi-core system >

➡ Reduce analysis running time!

Execution time of callbacks for the segment


$$R_{c,i}^{n+1} \leftarrow B_i + \sum_{\forall j: \tau_j \in \Phi_i} C_j + \sum_{\forall k: \tau_k \in e(\Phi_i) \vee \tau_k \in e_{HP}} \eta_i(R_{c,i}^n, \tau_k) \times C_k$$

Blocking time from lower priority callback

Interference from higher semantic priority chains

$$L_{\Gamma^c} = \sum_{\Phi_i \in \Gamma^c} R_{c,i}^n + S(\Gamma^c)$$

Blocking delay by prior instance



Evaluation

■ Experimental setup for case study

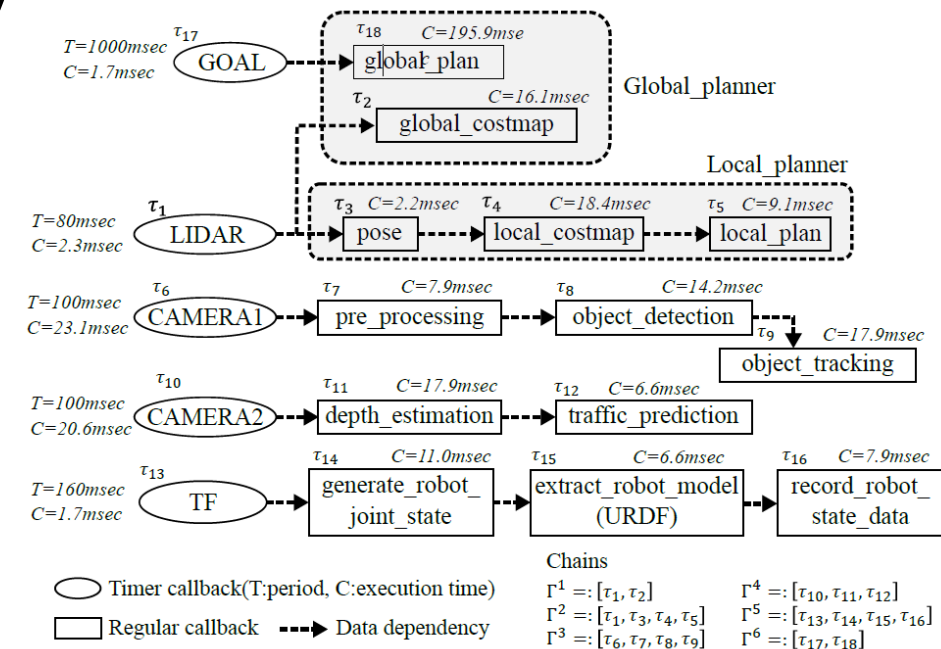
- Implemented in the Eloquent/Foxy/Galactic versions on Ubuntu18.04 on NVIDIA Xavier NX and Raspberry Pi 4.

■ Comparison of approaches

- ✓ **ROS2-SD[†]** : **ROS2 default scheduler** with resource reservation and **WCRT analysis** < F1TENTH >
- ✓ **ROS2-PiCAS** : proposed scheduler with end-to-end latency analysis

■ Case study in a multi-core system


- ✓ Inspired by the indoor self-driving stack of F1/10 vehicle
- ✓ 6 real-time chains (18 callbacks) and 6 best-efforts chains in a 4-core system
- ✓ Low-indexed chains are more critical chains



< Case study >

[†] D. Casini et al. "Response-time analysis of ROS 2 processing chains under reservation-based scheduling", ECRTS, 2019

Case study I




➡ **PiCAS schedules chains while respecting their semantic priorities.**

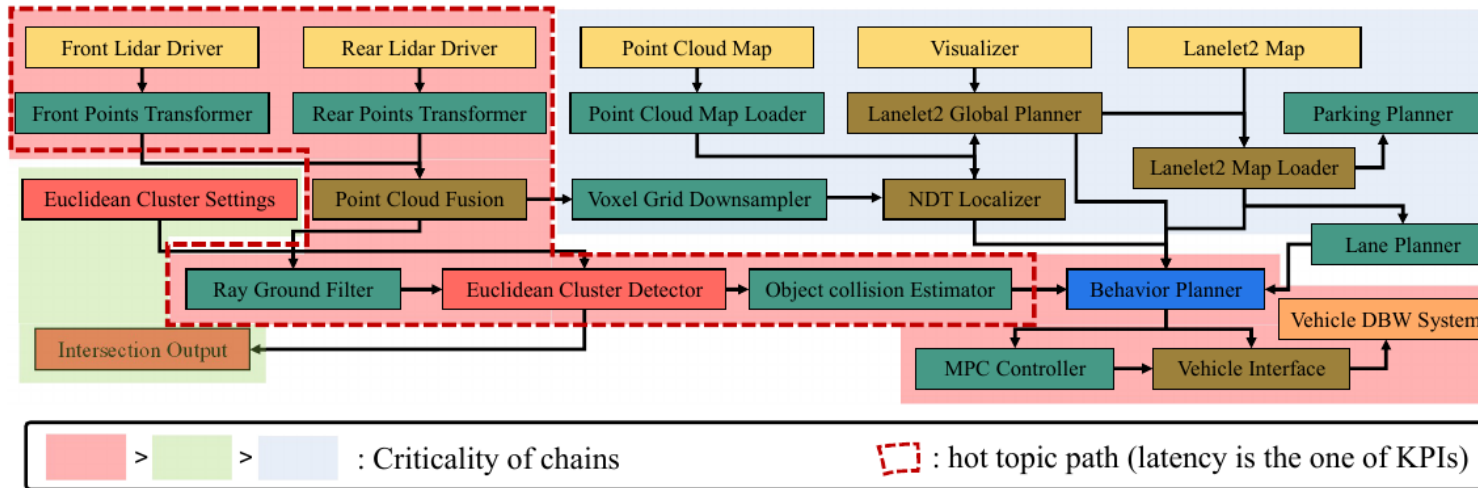
➡ **Our latency analysis provides tighter upper-bounds for real-time chains.**

Case study II

■ Integration to *Reference system*[†]



ROSCon 2021, Real-Time
Executor Workshop




[†] <https://github.com/rtenlab/reference-system>

< Chain configuration of Autware model >



< Average latency of hot topic path >



< Jitter of Behavior Planner >

Conclusion & Future work

■ Conclusion

- Proposed **a priority-driven chain-aware scheduling** and its **end-to-end latency analysis** framework
- New design of ROS2 scheduling includes **scheduling strategies**, **priority assignment of callbacks**, and **chain-aware node allocation**
- ROS2-PiCAS **outperforms** the existing ROS2 scheduling w.r.t. the end-to-end latency **under practical scenarios**

■ Future work

- Evaluate PiCAS to more complex scenario, e.g., autoware.auto (built on ROS2)
- Support multi-threaded executor of ROS2

Thank you

PiCAS: New Design of Priority-Driven Chain-Aware Scheduling for ROS2

- Hyunjong Choi, Yecheng Xiang, Hyoseung Kim

Q & A

PiCAS source: <https://github.com/rtenlab/ros2-picas>

PiCAS on reference system: <https://github.com/rtenlab/reference-system>

PiCAS paper: https://intra.ece.ucr.edu/~hyoseung/pdf/ras21_picas.pdf